

# LogisticRegression\_PimaIndians

September 19, 2020

## Coding Assignment 4

Logistic Regression Program - Pima Indians dataset

Author: Sreejith S

Ref: <https://www.kaggle.com/uciml/pima-indians-diabetes-database>

Ref: [https://www.medicinenet.com/glucose\\_tolerance\\_test/article.htm](https://www.medicinenet.com/glucose_tolerance_test/article.htm)

Ref: <https://towardsdatascience.com/understanding-confusion-matrix-a9ad42dcfd62>

```
[349]: # Importing the libraries
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

### 4.0. Load dataset

```
[350]: # Importing the dataset
dataset = pd.read_csv('/content/pima-indians-diabetes.csv')
print(dataset.shape)
```

(768, 9)

### 4.1. Understanding the data

#### 4.1.1. Inspecting Dataset

```
[351]: pd.set_option('display.max_columns', 10)
pd.set_option('display.width', 80)
print(dataset.head())
print("="*80)
print("="*80)
print(dataset.describe())
print("="*80)
print("="*80)
print(dataset.info())

#dataset.plot(kind='scatter', subplots=True, layout=(3, 3), sharey=False)
```

```
#plt.show()
```

	pregnant	plasma_glucose	dia_BP	skin_thickness	serum_insulin	bmi	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	diab_pedigree	age	Diab
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

	pregnant	plasma_glucose	dia_BP	skin_thickness	serum_insulin	\
count	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	
std	3.369578	31.972618	19.355807	15.952218	115.244002	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	

	bmi	diab_pedigree	age	Diab
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	pregnant	768 non-null	int64
1	plasma_glucose	768 non-null	int64
2	dia_BP	768 non-null	int64
3	skin_thickness	768 non-null	int64

```

4  serum_insulin    768 non-null    int64
5  bmi              768 non-null    float64
6  diab_pedigree    768 non-null    float64
7  age              768 non-null    int64
8  Diab             768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
None

```

```

[352]: #Creating info_dict for easier plotting
info_dict = {
    0:  ["pregnant", "# Pregnancies"],
    1:  ["plasma_glucose", "Plasma glucose concentration 2 hrs OGTT (mg/dl)"],
    2:  ["dia_BP", "Diastolic Blood Pressure (mm Hg)"],
    3:  ["skin_thickness", "Triceps skin fold thickness (mm)"],
    4:  ["serum_insulin", "2-Hour serum insulin (mu U/ml)"],
    5:  ["bmi", "BMI (kg/m^2)"],
    6:  ["diab_pedigree", "Diabetes pedigree function"],
    7:  ["age", "Age (years)"]
}

```

#### 4.1.2. Selecting class labels, and featureset

```

[353]: X = dataset.drop(columns=['Diab'])
print(X.head(), '\n')
y = dataset['Diab']
print(y.head())

```

	pregnant	plasma_glucose	dia_BP	skin_thickness	serum_insulin	bmi	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	diab_pedigree	age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33

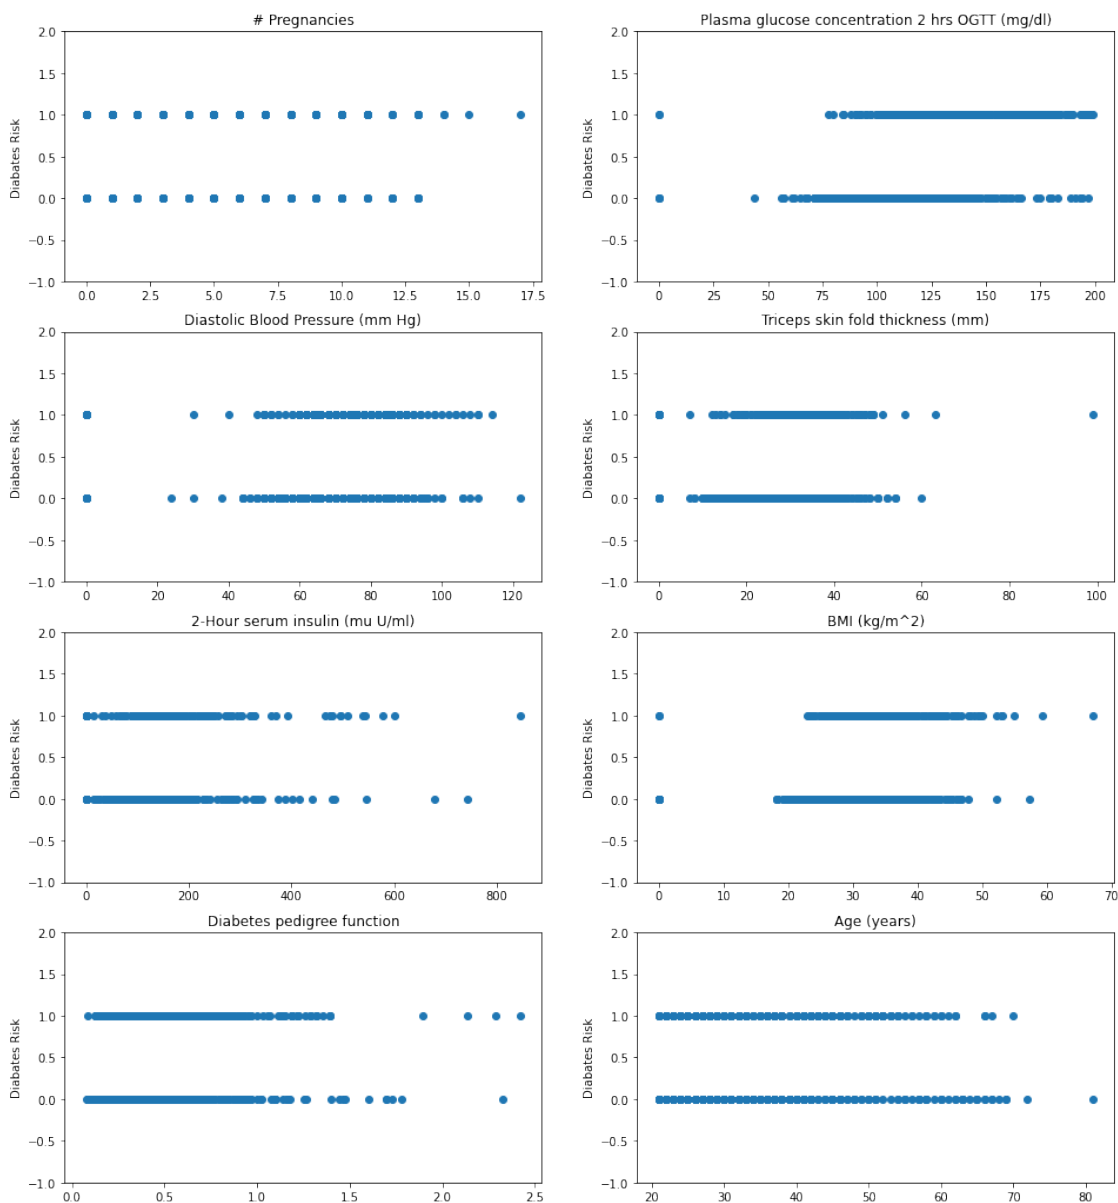
0	1
1	0
2	1
3	0
4	1

Name: Diab, dtype: int64

### 4.1.3. Plotting the dataset

```
[354]: fig = plt.figure(figsize=(16,18))
for i in range(8):
    ax = fig.add_subplot(4, 2, i+1)
    ax.scatter(X.iloc[:, i], y)
    ax.set_title(info_dict[i][1])
    ax.set_ylabel("Diabetes Risk")
    ax.set_ylim(-1,2)

plt.show()
```



#### 4.1.4. Observations

**Invalid data points** It can be observed that several features have zero values when it is impossible to have zero values for them. plasma\_glucose, dia\_BP, BMI etc are a few such features. Either these outliers can be removed from the dataset or can be replaced with mean values.

**Which input variable, do you think, is better in predicting if someone is going to get diabetes in the next 5 year?** It is hard from the plots in 4.1.3 alone to predict which input variable is better in predicting if someone is going to get diabetes in the next 5 years. On visual inspection, BMI, plasma\_glucose seems to have some effect on the outcome.

Comparing means of each feature between diabetic and non-diabetic might give some insights. But the means can be affected by the outliers. From the code cell below, it can be observed that diabetic patients have a higher BMI=35 as compared to BMI=30 for non-diabetics.

Another option might be to see the correlation between features and Diab. Padas provide a method corr() for dataframes. This was tested in the cell below and BMI, plasma\_glucose, diab\_pedigree, serum\_insulin & even age to some extent was found to have an influence on diabetes.

```
[355]: for feature in X:
        m1 = np.mean(dataset[dataset.Diab == 1][feature])
        m0 = np.mean(dataset[dataset.Diab == 0][feature])
        print(f"Avg value of {feature:<16} when, y = 1, is {m1}")
        print(f"Avg value of {feature:<16} when, y = 0, is {m0}")
        print("\n")

dataset.corr(method='pearson')
```

Avg value of pregnant	when, y = 1, is 4.865671641791045
Avg value of pregnant	when, y = 0, is 3.298
Avg value of plasma_glucose	when, y = 1, is 141.25746268656715
Avg value of plasma_glucose	when, y = 0, is 109.98
Avg value of dia_BP	when, y = 1, is 70.82462686567165
Avg value of dia_BP	when, y = 0, is 68.184
Avg value of skin_thickness	when, y = 1, is 22.16417910447761
Avg value of skin_thickness	when, y = 0, is 19.664
Avg value of serum_insulin	when, y = 1, is 100.33582089552239

Avg value of serum\_insulin      when, y = 0, is 68.792

Avg value of bmi                      when, y = 1, is 35.14253731343278

Avg value of bmi                      when, y = 0, is 30.304199999999996

Avg value of diab\_pedigree      when, y = 1, is 0.5505

Avg value of diab\_pedigree      when, y = 0, is 0.429734000000000017

Avg value of age                      when, y = 1, is 37.06716417910448

Avg value of age                      when, y = 0, is 31.19

```
[355]:
```

	pregnant	plasma_glucose	dia_BP	skin_thickness	\
pregnant	1.000000	0.129459	0.141282	-0.081672	
plasma_glucose	0.129459	1.000000	0.152590	0.057328	
dia_BP	0.141282	0.152590	1.000000	0.207371	
skin_thickness	-0.081672	0.057328	0.207371	1.000000	
serum_insulin	-0.073535	0.331357	0.088933	0.436783	
bmi	0.017683	0.221071	0.281805	0.392573	
diab_pedigree	-0.033523	0.137337	0.041265	0.183928	
age	0.544341	0.263514	0.239528	-0.113970	
Diab	0.221898	0.466581	0.065068	0.074752	

	serum_insulin	bmi	diab_pedigree	age	Diab
pregnant	-0.073535	0.017683	-0.033523	0.544341	0.221898
plasma_glucose	0.331357	0.221071	0.137337	0.263514	0.466581
dia_BP	0.088933	0.281805	0.041265	0.239528	0.065068
skin_thickness	0.436783	0.392573	0.183928	-0.113970	0.074752
serum_insulin	1.000000	0.197859	0.185071	-0.042163	0.130548
bmi	0.197859	1.000000	0.140647	0.036242	0.292695
diab_pedigree	0.185071	0.140647	1.000000	0.033561	0.173844
age	-0.042163	0.036242	0.033561	1.000000	0.238356
Diab	0.130548	0.292695	0.173844	0.238356	1.000000

#### 4.1.5. Clean-up dataset

Fixing zero values - On exploring the kaggle discussions on the dataset, it was found that the zero values are due to incomplete hospital records. We can compare effects on replacing these zero values.

```
[356]: replace_zeros_with_mean = False
if replace_zeros_with_mean:
    #testing
    print(dataset[dataset.bmi == 0])
```

```

print("-"*80)

cols = ["bmi", "dia_BP", "skin_thickness", "plasma_glucose"]
means = []
for c in cols:
    mean = np.mean(dataset[dataset[c] != 0][c])
    means.append(mean)

for c, m in zip(cols, means):
    dataset[c] = dataset[c].replace(0, m)

#testing
print(dataset[dataset.bmi == 0])

```

## 4.2. Write code to find out how many people in the dataset are marked as having 5-year diabetes?

From the code cell below, 268 patients are marked as having diabetes.

```

[357]: dataset[(dataset.Diab == 1)]
#print(dataset[dataset.Diab == 1].shape[0])

```

```

[357]:
    pregnant  plasma_glucose  dia_BP  skin_thickness  serum_insulin  bmi  \
0           6           148      72           35           0  33.6
2           8           183      64           0           0  23.3
4           0           137      40           35          168  43.1
6           3           78      50           32           88  31.0
8           2          197      70           45          543  30.5
..         ...           ...      ...           ...           ...  ...
755          1          128      88           39          110  36.5
757          0          123      72           0           0  36.3
759          6          190      92           0           0  35.5
761          9          170      74           31           0  44.0
766          1          126      60           0           0  30.1

    diab_pedigree  age  Diab
0           0.627   50     1
2           0.672   32     1
4           2.288   33     1
6           0.248   26     1
8           0.158   53     1
..           ...   ...   ...
755          1.057   37     1
757          0.258   52     1
759          0.278   66     1
761          0.403   43     1
766          0.349   47     1

```

[268 rows x 9 columns]

### 4.3. Creating Logistic Reg. model and visualizing

#### 4.3.1. Creating feature sets

There separate feature sets are created as per Q4.3

```
X <- All features
```

```
X_4_3_a <- 'age', 'pregnant', 'bmi'
```

```
X_4_3_b <- 'skin_thickness', 'diab_pedigree', 'dia_BP'
```

```
[358]: X = dataset.drop(columns=['Diab'])
X_4_3_a = dataset[['age', 'pregnant', 'bmi']]
X_4_3_b = dataset[['skin_thickness', 'diab_pedigree', 'dia_BP']]
print(X.head(), '\n')
print("-"*80)
print(X_4_3_a.head(), '\n')
print("-"*80)
print(X_4_3_b.head(), '\n')
print("-"*80)
y = dataset['Diab']
print(y.head())
```

	pregnant	plasma_glucose	dia_BP	skin_thickness	serum_insulin	bmi	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	diab_pedigree	age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33

---

	age	pregnant	bmi
0	50	6	33.6
1	31	1	26.6
2	32	8	23.3
3	21	1	28.1
4	33	0	43.1

---



	skin_thickness	diab_pedigree	dia_BP
0	35	0.627	72
1	29	0.351	66
2	0	0.672	64
3	23	0.167	66
4	35	2.288	40

```
-----
0    1
1    0
2    1
3    0
4    1
```

Name: Diab, dtype: int64

### 4.3.2. Logistic Regression

#### Using all 8 input features

```
[359]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split

def create_model(X, y, test_size=0.3, random_state=0, quiet=False):
    """
    """
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size,
→random_state=random_state)
    logreg = LogisticRegression(max_iter=200)
    logreg.fit(X_train, y_train)

    y_pred = logreg.predict(X_test)
    test_acc = logreg.score(X_test, y_test)
    train_acc = logreg.score(X_train, y_train)
    if not quiet:
        print("Accuracy of Logistic Regression Classifier: ")
        print(f'on test set: {test_acc:.2f}')
        print(f'on train set: {train_acc:.2f}')
        print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
    influence = np.std(X, 0)*logreg.coef_[0]
    return test_acc, train_acc, influence

_, _, influence1 = create_model(X, y)
```

Accuracy of Logistic Regression Classifier:

on test set: 0.78

on train set: 0.77

Confusion Matrix:

```
[[141  16]
 [ 35  39]]
```

#### a. Only age, pregnant, and bmi as the input features

```
[360]: create_model(X_4_3_a, y)
```

Accuracy of Logistic Regression Classifier:

on test set: 0.68

on train set: 0.69

Confusion Matrix:

```
[[135  22]
 [ 51  23]]
```

```
[360]: (0.683982683982684, 0.6871508379888268, age          0.488079
pregnant      0.218059
bmi           0.869156
dtype: float64)
```

#### b. Only skin\_thickness, diab\_pedigree, and dia\_BP as the input features

```
[361]: create_model(X_4_3_b, y)
```

Accuracy of Logistic Regression Classifier:

on test set: 0.68

on train set: 0.64

Confusion Matrix:

```
[[151   6]
 [ 68   6]]
```

```
[361]: (0.6796536796536796, 0.6443202979515829, skin_thickness  0.134746
diab_pedigree      0.276704
dia_BP             0.136894
dtype: float64)
```

'plasma\_glucose', 'serum\_insulin', 'bmi', 'diab\_pedigree'

```
[362]: #picked from top correlations
X_4_3_x = dataset[['plasma_glucose', 'serum_insulin', 'bmi', 'diab_pedigree']]
_, _, influence2 = create_model(X_4_3_x, y)
```

Accuracy of Logistic Regression Classifier:

on test set: 0.80

on train set: 0.76

Confusion Matrix:

```
[[143  14]
 [ 33  41]]
```

**4.4. Out of the three models (all features, and then (a), (b) features), which model performs the best in classification? Specify the reason too.**

**Without Replacing outlier zeros**

Features Selected	Test Accuracy	Train Accuracy	Confusion Matrix (tn, fp, fn, tp)
all features	0.78	0.77	141, 16, 35, 39
4.3.a) age, pregnant, and bmi	0.68	0.69	135, 22, 51, 23
4.3.b) skin_thickness, diab_pedigree, dia_BP	0.68	0.64	151, 6, 68, 6
plasma_glucose, serum_insulin, bmi, diab_pedigree	0.80	0.76	143, 14, 33, 41

**After Replacing Outlier Zeros with Mean**

Features Selected	Test Accuracy	Train Accuracy	Confusion Matrix (tn, fp, fn, tp)
all features	0.77	0.77	141, 16, 36, 38
4.3.a) age, pregnant, and bmi	0.69	0.69	136, 21, 50, 24
4.3.b) skin_thickness, diab_pedigree, dia_BP	0.72	0.65	150, 7, 57, 17
plasma_glucose, serum_insulin, bmi, diab_pedigree	0.79	0.77	140, 17, 31, 43

The model with all 8 features performs best in the original three cases. This was because the other two 4.3.a & 4.3.b might not have been able to capture the true model as soem significant feature, like plasma\_glucose, might have been omitted.

On exploring to find if a better feature set existed it was found that (plasma\_glucose, serum\_insulin, bmi, diab\_pedigree) performed better than the 8-input feature set. For such a small dataset, it might be hard to know if this will always hold true. But finding conforms with existing literature on diabetes risks.

Replacing zeros with mean didn't improve the best performing sets. Perhaps, this might be due to the fact that the outliers are preventing the model from overfitting to any particular feature too much.

Interestingly, 4.3.b) skin\_thickness, diab\_pedigree, dia\_BP had the best true-negative among all, but it performed poorly on every other confusion matrix values (fn, fp & tp).

It was also observed that, using a test\_size=0.2 also improves the test accuracy.

Since we were able to achieve a better performance with fewer input features (plasma\_glucose, serum\_insulin, bmi, diab\_pedigree), the remaining input features can be considered insignificant for this particular dataset. It was also noted that the aforementioned four features had a higher correlation with Diab, when compared to features like skin\_thickness.

In the cell below, the influence of each feature can be observed from the `logreg.coef_` which are the magnitude of the coefficients, multiplied by S.D of each feature set [Ref](#). Even though, dia\_BP has a higher magnitude than serum\_insulin when all features were used, it doesn't improve performance - [code in last cell]

```
[363]: print("8 input feature set")
print(influence1.sort_values(axis=0))
print('-'*40)
print("4 input feature set")
print(influence2.sort_values(axis=0))
```

```
8 input feature set
dia_BP      -0.226341
serum_insulin -0.145447
skin_thickness 0.099055
diab_pedigree 0.226526
pregnant     0.290324
age          0.292800
bmi          0.708817
plasma_glucose 1.089174
dtype: float64
-----
4 input feature set
serum_insulin -0.195370
diab_pedigree 0.230057
bmi          0.623056
plasma_glucose 1.181423
dtype: float64
```

```
[366]: #Trying to find if a better accuracy feature set exist iteratively since we have
→ a small dataset
from itertools import combinations
features = ["pregnant", "plasma_glucose", 'dia_BP', 'skin_thickness',
→ 'serum_insulin', 'bmi', 'diab_pedigree', 'age']
combs = list(combinations(features, 4)) #get all combinations of r=4
#sum([list(map(list, combinations(features, i))) for i in range(1, len(features)
→ + 1)], [])
best_test = 0
best_feature_set = []
for c in combs:
    X_4_3_x = dataset[list(c)]
    test_acc, _, _ = create_model(X_4_3_x, y, quiet=True)
    if test_acc > best_test:
```

```
best_test = test_acc
best_feature_set = c

print(f"best 4 input feature set - {best_feature_set}, test_acc = {best_test:.
→2f}")
```

```
best 4 input feature set - ('plasma_glucose', 'serum_insulin', 'bmi',
'diab_pedigree'), test_acc = 0.80
```